

CS 470/548: Assignment 05

Programming Assignments (100% undergraduate, 80% graduate)

The goal of this assignment is to write code to **train and evaluate PyTorch neural network models to predict classes from the CIFAR10 dataset.**

CS 470: You must train and evaluate at least TWO different variations.

CS 548: You must train and evaluate at least FOUR different variations.

A05.py

You will write the following functions in this python script:

- `get_approach_names()`
- `get_approach_description(approach_name)`
- `get_data_transform(approach_name, training)`
- `get_batch_size(approach_name)`
- `create_model(approach_name, class_cnt)`
- `train_model(approach_name, model, device, train_dataloader, test_dataloader)`

Note that most of these functions rely on the `approach_name` parameter. This is how you will code different behavior per approach.

`get_approach_names()`

Returns a list of the names of all combinations you will be testing. It is largely up to you what names you use, but self-documenting names are encouraged.

A rather boring example: ["CNN0", "CNN1"]

`get_approach_description(approach_name)`

Given the `approach_name`, return a text description of what makes this approach distinct. This does not have to be very long: one sentence is sufficient.

`get_data_transform(approach_name, training)`

Given the `approach_name` and whether this is for training data, return the appropriate dataset transform. You may add data augmentation if you like for training data transforms, BUT:

- Do NOT do data augmentation for non-training data!
- At minimum, you must perform the following:

```
data_transform = v2.Compose([v2.ToImageTensor(), v2.ConvertImageDtype()])
```

`get_batch_size(approach_name)`

Given the `approach_name`, return the preferred batch size. This can be hardcoded to a single value if you do not have a preference, BUT this allows you to use a smaller batch size if your architecture is large enough where that is a concern.

`create_model(approach_name, class_cnt)`

Given the `approach_name` and output `class_cnt`, build and return a PyTorch neural network that takes a color 32 x 32 image and outputs a vector with `class_cnt` elements.

You do not have to move this model to the GPU (that is done for you in the provided scripts).

Any of your networks should have the following attributes:

1. The input channels should be 3.
2. The output layer should have `class_cnt` nodes.
3. The network must have at least 3 convolutional and/or fully-connected layers, not including the last output layer.
4. **You are free to use ANY of the layer types available in PyTorch.**
5. You may choose to use one of the existing pre-trained networks (like VGGNet) and fine-tune it.

`train_model(approach_name, model, device, train_dataloader, test_dataloader)`

Given the provided model, the device it is located, and the relevant dataloaders, train this model and return it.

A couple of caveats:

- **You can ONLY train on the data from `train_dataloader`.**
- You may print out how the network is doing on `test_dataloader`. **However, you may NOT use `test_dataloader` as a validation set** (e.g., you cannot use it to tune learning rates dynamically).
- You are free to choose any optimizer you wish.
- You are free to decide how many epochs of training you use.
- **Your loss SHOULD be `CrossEntropyLoss`.** You are, however, welcome to add other losses if you believe it will help.

Note that, in general, differences here do NOT count as acceptable approach differences, such as:

- **Epoch counts**
- **Different optimizers**

Acceptable Approach Differences

Your approaches should be sufficiently different from each other! Acceptable differences include:

- Different sizes of filters
- Notably different filter counts (32 vs. 64, not 32 vs. 33)
- Notably different node counts
- Consistently different activation functions.
- Adding 1 or more additional layers
- Adding skip connections
- With or without significant data augmentation (or sufficiently different variations of data augmentation)

I actively encourage you to look for examples of working networks.

HOWEVER, if you do use code from another source, cite it in the body of the code like so:

Taken from: <https://keras.io/getting-started/sequential-model-guide/>

NOTE: Even if you cite it, you may NOT use code from another student (whether presently in the course or from a previous semester).

However, any code in the slides for this course is freely usable.

If you are unsure whether any variants you intend to experiment with are different enough, ASK ME.

TRAINING CAN BE VERY TIME-CONSUMING. DO NOT LEAVE THIS UNTIL LAST MINUTE!

Provided Programs

Train_A05.py

- Asks you which approach you wish to train.
- Loads the data.
- Trains the model.
- Saves the model to assign05/output

Eval_A05.py

- Asks you which approach you wish to evaluate (if -1, evaluates all).
- Evaluates the model(s) on training and testing data.
- Prints and saves the results for each approach to assign05/output with filename: <approach name>_RESULTS.txt or ALL_RESULTS.txt.

CS 548 ONLY: Report (20%)

In addition to the code, you must **submit a report** that contains the following:

- CS 548, Assignment 05, and your name
- A **brief summary** of your approaches and the differences between them
- A **brief discussion** as to WHY you thought these variations would have made a difference
- A **table** of results for each network, showing:
 - o **Training accuracy**
 - o **Training F1 score**
 - o **Testing accuracy**
 - o **Testing F1 score**
- Any **citations** to external network architectures used (any reasonable citation method is allowed).

BONUS POINTS

For the entire assignment, bonus points will be awarded to the submission that has the following:

- (+%5) Best **testing accuracy** in class
- (+%5) Best **testing F1 score** in the class
- (+%5) Experimenting with N variants or more
 - o **CS 470: N = 4**
 - o **CS 548: N = 8**

"Best in class" will be assessed separately for undergraduate and graduate students.

Submission

Make sure the following is committed on your repo:

- **A05.py**
- The contents of **assign05/output**, including:
 - o **Your model files (model_*.pth) UNLESS they are over 10MB!**
 - o **ALL_RESULTS.txt**
- **CS 548 only: Your report as a PDF**

PLEASE add the following to your .gitignore file:

- data/
 - o This is where CIFAR10 is downloaded
- assign05/output/bigmodel.pth
 - o Any models over 10MB

Grading

CS 470: 100% Programming

CS 548: 80% Programming, 20% Report